

AI in Education: How Students, Teachers, and Schools Are Responding

A Research Report For Educators Based on Six International Schools in Switzerland

Niklas Obermüller and Hiroshi Tokunaga

Williams College Global Scholars Programme, January 2026

April 20th, 2026

Table of Contents

AI in Education: How Students, Teachers, and Schools Are Responding	1
Table of Contents	2
I. Introduction	3
What is this report?	3
Who are we?	3
Why did we conduct this research?	4
Methodology	4
II. School Landscape	5
What are the different approaches adopted by schools?	5
III. How Are Students and Teachers Actually Using AI?	7
How widespread is AI use?	7
What are students and teachers using AI for?	7
How are different groups of students using AI?	8
How are different groups of teachers responding to AI?	11
IV. What Is AI Doing to Learning, Teaching, and Assessment?	13
Do students learn better or worse with AI?	13
How does AI change the way teachers teach?	13
What is happening to assessment?	14
V. What Do Students and Teachers Actually Want?	15
Why are students using AI?	15
What do students want from their schools?	15
What do teachers want from their schools?	15
VI. Does School Policy Produce the Intended Outcomes?	17
What do students and teachers think about their AI policies?	17
Do students and teachers comply?	17
Do teachers actually know how students are using AI?	18
VII. What Can Schools Control, and What Can't They?	20
What sits outside schools' direct control?	20
What does the evidence suggest schools should prioritise?	20
Appendix A: Interview Questions	23
Appendix B: Notes	26

I. Introduction

What is this report?

In January 2026, two student researchers conducted research on AI in Education at six international schools in Switzerland. This report is one of the outcomes of the research, written specifically for school educators. The central question this report explores is how different stakeholders of education respond to the impact of AI. By “AI,” we specifically mean generative AI, the wave of tools that took off after ChatGPT was released to the public in late 2022. We organized our findings into five sections that address the questions we think matter most to educators:

1. **How are students and teachers actually using AI?** What they use it for, how often, and how use differs across groups.
2. **What is AI doing to learning, teaching, and assessment?** Whether students learn better or worse, how assessment is changing, and what teachers see in their classrooms.
3. **What do students and teachers actually want?** Why students turn to AI, what they want from schools, and what teachers are asking for.
4. **Does school policy produce the intended outcomes?** What policies exist, whether people follow them, and where perception gaps emerge.
5. **What can schools control, and what can’t they?** The structural constraints schools face and where the evidence suggests they should focus on.

This is a general report that examines trends across the six schools studied, so findings come from aggregated data. We anonymize school names and represent them by letters (e.g. School A); we also remove identifying characteristics when using interviewee quotes. School-specific data is delivered to the schools via a separate document.

Who are we?

We are Niklas Obermüller and Hiroshi Tokunaga. We are current second-year students at Williams College, USA. Hiroshi graduated from Aiglon College in 2024. Both Niklas and Hiroshi were interested in AI’s impact on education, and were selected as Global Scholars. The Global Scholars program chose Geneva as the destination for the January research project. Hiroshi reached out to Swiss educator Mr Richard McDonald, and was in the middle of planning a project with him before he tragically passed away. In the spirit of continuing Richard’s legacy, we turned to Richard’s close friend and collaborator, Mr Ilya Eigenbrot, for guidance. Ilya generously helped us coordinate with many of the schools. We dedicate this work in memory of Richard.

We were both motivated to do this research because we see great potential and uncertainties in how AI is changing the educational landscape. Our hope is that through this research, we can contribute to a shared body of knowledge about the impact of AI on education.

Why did we conduct this research?

We think AI is having an enormous impact on education as a system, so rather than study each school as individual institutions we looked for patterns across schools, and included data from these stakeholders: administrators, teachers, students. We aspire to gather data from policymakers and parents in future studies, as they are important stakeholders we did not manage to speak to.

While the schools studied represent a small sample of private institutions in Switzerland, we believe they offer insights into how AI is impacting education on a broader scale. These are some of the most well-resourced schools in the world, with diverse students and the capacity to adapt faster than public systems. Most existing research on AI in education focuses on public school systems. If even these schools are finding it challenging to navigate AI, it shows something important about the challenges facing education everywhere.

Methodology

We collected data in January 2026 through a mixed-methods approach: surveys first, then interviews.

Surveys. We designed two surveys, one for students and one for teachers. 276 students and 149 teachers across five schools completed the surveys before our school visits. One school participated only in the interview phase. The surveys asked about how often people use AI, what they use it for, how well they know their school's AI rules, and what they think about AI's effects on learning and assessment.

Interviews. We then visited all six schools and conducted semi-structured interviews with administrators, teachers, and students on-site. Administrators typically include the Head of School, Head of Academics and the Head of AI/Technology. We spoke with roughly four teachers per school, selected to represent different subjects: STEM, languages, arts, and humanities. We also interviewed students from a range of year groups, typically grade 8 to 12. Interviews with administrators and teachers lasted 20 to 40 minutes each, while those for students lasted about 10 to 15 minutes each.

Analysis. All school names were anonymized during the analysis of surveys and interviews. We analyzed survey responses using descriptive statistics and compared across subgroups such as graduation year, usage frequency, subject, and teaching experience. We analyzed interviews thematically, looking for patterns that appeared across schools. Throughout this report, we are careful to flag when a finding comes from a small sample or a single school.

Limitations. Students may underreport their AI use even in anonymous surveys. Some schools' safeguarding policy required a teacher in the room when we interviewed students, so this could bias student interview data. Also, some teacher and student subgroups at certain schools are very small, so results stemming from their data should be seen as trends rather than provable findings. The sample is purposive, not random: we studied these six schools because we could access them, so findings should be generalized with caution for readers outside of these institutions.

II. School Landscape

We interviewed administrators at six schools. What follows draws directly from those conversations. If you are thinking about your school's AI strategy, we hope hearing from peer institutions navigating the same terrain can help you make more informed decisions.

The six schools reacted to generative AI at very different speeds. Some formed working groups within months of ChatGPT's release; others are still developing their first formal policy as of January 2026. But there was one thing every administrator was adamant about: AI has not changed what their school is for. Not a single administrator described AI as requiring a revision of their institutional mission or values. Instead, AI has encouraged them to rethink the approaches to achieving their values.

What are the different approaches adopted by schools?

Approach 1: Fast, systematized response to AI. Within months of ChatGPT's release in late 2022, they formed a dedicated AI working group and over two years constructed the most comprehensive AI governance model we encountered: a colour-coded 1-to-5 assessment scale, a position paper, document-tracking, and three specialised working subgroups focusing on different issues. The director described that the school succeeded in creating AI literacy across the community, but is still in the process of addressing the hardest question, which is how to preserve the cognitive struggle in learning. They are now in the stage of converting the position paper into a formal policy.

Approach 2: Have a dedicated human role and use change management. This school has a full-time Head of AI position and built its entire strategy around it. The Head of AI researches and filters tools, designs and delivers training, runs pilot courses, and manages a traffic-light system (red/amber/green) embedded in the learning management system. By January 2026, the school was piloting an AI course for Year 10 students. The general director's philosophy is drawn from change management theory: you cannot convert the 15 to 20% of teachers who resist, and you should not focus on them first. The priority is bringing the majority in line with the policies and implementations. The school prefers having "guidelines" over "policies" because guidelines encourage teachers to actively implement them.

Approach 3: Let teachers make the rules and values guide the use of AI. Two teachers wrote the AI charter and introduced it through homeroom and computer science classes. AI café is hosted every month for teachers to exchange ideas and ask questions. The head wants every student to ask themselves one question before using AI: "Am I using this tool responsibly?". The school's core values of kindness, respect, and responsibility are described as "as valid in 2026 as they were in 1826". For them, AI does not require new values, just a more careful application of existing ones. To back this up practically, the school shifts more assessment into class time and increases the weight of oral exams.

Approach 4: Design the policy around learning pedagogy and teachers. This school chose to focus on what AI could be used for in each subject, rather than which tools are

permitted or banned. The approach is pedagogically ambitious but is still in development as of January 2026. Full implementation was projected at six to nine months. They identified a sequencing problem: before you can give teachers guidance on how students should use AI, you need to ensure teachers themselves feel confident with it, and that baseline training is still underway.

Approach 5: Trust the shared understanding that exists between teachers. This school has no written AI policy yet, but teachers talk to each other about the norms informally. When a policy does come, data protection and keeping students safe will be its top priority. The school runs on Google, so students use Gemini rather than ChatGPT; data stays within a school-managed environment. Younger students in junior school have no AI access at all. The director says administrators can't really know what students are doing with AI — it is like seeing the tip of an iceberg.

Approach 6: Adopt the IB's recommended framework and build from there. This school has structured its AI rules around the IB's official guidance, with a classroom poster mapping scenarios to acceptable and not acceptable use. The school's orientation toward AI is optimistic, and is exploring potential opportunities for AI to enhance and innovate learning experiences. Google Gemini is available for students in the senior school.

III. How Are Students and Teachers Actually Using AI?

How widespread is AI use?

AI use among students is nearly universal. Only 6% of the 276 surveyed students report never using AI during the school week. 42% report using it once or twice per week, 28% report using it 3 to 4 times, and 24% say they use it 5 or more times. For most students, AI appears to be a routine part of the school week rather than something they are cautiously experimenting with.

How often do students use AI during the school week?

N = 276 students across five schools




Figure 1. Student AI use frequency

Among teachers, 66% report using AI for planning, materials creation, or grading. But the distribution is sharply polarised: 22% actively disagree with using AI for professional tasks. There is little middle ground. **Teachers have either adopted AI or they have not, and 59% of those who have adopted it agree it has improved their efficiency.**

In interviews, a small subset of students described using multiple AI platforms simultaneously, including NotebookLM, Perplexity, and Grok alongside ChatGPT and Gemini. This suggests a developing AI literacy among heavier users that extends beyond any single tool. Among teachers, a meaningful minority of non-adopters cite principled concerns about cognition, creativity, and professional identity rather than unfamiliarity. One teacher described AI-generated report writing as making him feel morally uncomfortable; another experienced social pressure from academic colleagues for openly referencing AI in a departmental meeting.

What are students and teachers using AI for?

The most common student tasks are checking answers (48%), brainstorming ideas (46%), summarising texts (39%), solving maths and science problems (38%), and translation or language support (37%). Direct drafting of written assignments is reported by 23% of students. Yet, cross-referencing this with interview data suggests this number substantially underestimates actual practice. The least common use was creating study plans (21%), a pattern consistent across all six schools independently.




Figure 2. Student AI use cases, multi-select

Among teachers, AI use concentrates on planning and materials creation. Multiple teachers name translation as their single most impactful AI application. Translation is not only used by language teachers: teachers translate materials for non-native English speakers to help them better understand the material. Teachers draw a consistent line around using AI to evaluate student work: at School B, for example, 62% of teachers disagree or strongly disagree with using AI for student assessment. 86% report being transparent with students about their own AI use and 73% say they actively guide students on AI best practices.

In interviews, students make an implicit distinction between **AI as a thinking partner** and **AI as a task completer**. Some students use AI to augment learning, such as generating revision material from their own notes or creating mock exams from their own study sheets. Others describe mixed use, such as brainstorming with AI and generating inspiration for drafts as a starting point for their own writing. And some describe clear substitution, such as asking AI to complete a full written assignment or requesting solutions without engaging with the material.

On the teacher side, a small number have moved beyond standard use into what might be called **pedagogical redesign**: building custom AI agents constrained to asking questions rather than supplying answers, or designing projects that structurally cannot be completed by AI alone. This approach, making tasks AI-proof by design, is fundamentally different from detection or prohibition.

How are different groups of students using AI?

Older students use AI more intensively. A clear gradient runs from the graduating class to the youngest cohorts. 39% of 2026 graduates use AI 5 or more times per week, then declining steadily among early secondary students.




Figure 3. Student AI use frequency by expected graduation year

The 2026 cohort, the heaviest users in the sample, also scores highest on “AI harms my understanding” (4.37 vs 3.92 sample mean). These students use AI the most and worry most about its cognitive effects. This pattern is consistent with what we heard in interviews: **examination pressure drives AI adoption even among students who recognise its risks.** The “later/unknown” cohort, predominantly early secondary students, scores lowest on “AI helps me learn” (4.78 vs 5.53) and shows the highest non-use rate. One contributing factor for this is that younger students are restricted in using their devices.




Figure 4. Student attitudes by expected graduation year

Language background shapes how AI is used. Native English speakers are heavier users: 32% use AI 5 or more times per week, compared to 19% for non-native students. When asked to consider the statement “For certain tasks, I rely on AI more than on direct guidance from a teacher.” non-native students agree more than native students (4.18 vs. 3.93 means). Non-native students report that they use it mainly for language support and translation and 32% of them use AI primarily in a language other than English. As international students ourselves, we think that a potential explanation is that students with English not as their native language use AI when they face a language barrier, especially for tasks that require jargon in English or more complex language.

High achieving students use AI more. Students whose self-reported academic standing are below average (N=11) show the highest belief that AI helps learning (90.9%), yet the lowest AI use intensity. This suggests that students who feel most academically have not yet integrated AI as a regular learning tool. Students who self-reported to be above-average, by contrast, demonstrate both high use and high rule-override, meaning that perceived academic competence correlates with strategic (if sometimes non-compliant) AI use.

High usage frequency in students correlates with other survey items. Heavy users of AI report greater reliance on AI over teacher guidance, more willingness to use AI when assignments discourage it (+2.07), stronger belief that AI helps them learn (+1.88), more confidence in its output (+1.39), and lower compliance with teacher-set AI rules (−1.05). Whether this reflects AI use gradually shaping student attitudes, or the reverse — students who already trusted AI and were comfortable bending rules simply using it more — cannot be determined from our data. Confounding is a third possibility: academically struggling students, or those in classrooms with weaker AI guidance, may both use AI more and hold these attitudes for reasons unrelated to AI use itself. Self-report consistency bias also cannot be ruled out, since students who characterized themselves as heavy users may have felt motivated to justify that use by rating AI favorably, artificially inflating the observed correlations.




Figure 5. Student survey responses by AI use frequency

How are different groups of teachers responding to AI?

Associations between subject area, AI use, and AI concerns. In the surveys, we asked teachers to select the subjects they teach, then grouped the subjects into four categories: STEM (Maths & Sciences), Humanities, Languages, Arts.

The language teacher group, which primarily consists of second-language instructors but also includes some English literature teachers, reports the highest level of AI use for planning (mean 3.83 vs 3.50). They also informed us of a structural advantage in detection: known vocabulary and grammar patterns at each proficiency level make AI-generated language output easier to identify.

STEM (Maths & Sciences) teachers report the lowest level of AI use for planning lessons. However, interviews reveal some of the most sophisticated AI use in the dataset, including prompting agents to a certain behavior like providing explanations rather than solutions, and workflow-embedded generation.

Humanities teachers show the highest concern about AI's effect on assessment, the highest detection difficulty, and the lowest support for AI integration into education (mean 3.05 vs 3.49). This makes structural sense: assessment in the humanities relies on extended written output, which is easy to generate with AI.




Figure 6. Teacher AI engagement by subject area

Amount of teaching experience is associated with differences in AI attitudes and classroom practices. Teachers with 30+ years of experience use AI less than all other groups (3.00 vs. 3.50), they adapted the AI policy least for their classes (2.44 vs. 2.86). They agree least with AI being a useful learning tool. What is also interesting is that they are less alarmed than their colleagues: less convinced of student misuse (3.44 vs 3.95), less worried about assessment impact (3.25 vs 3.81), and least reliant on detection tools (2.12 vs 2.68). By contrast, teachers just starting their career, with 3 or less years of teaching, place consistently on the other side of the spectrum.

Teacher policy familiarity. From the graph below, we can see that teachers who self-report to be more familiar with AI policy are generally more confident in using AI for their teaching, provide students feedback on their AI use more regularly, and implement the schools policy more consistently. Outside of the graph, the survey also reveals that teachers unfamiliar with the policy are more convinced that students are transparent about their AI use, and they are less concerned that AI complicates accurate student assessment (2.92 vs.

3.81). An explanation is that their positivity may reflect not seeing the problem rather than informed enthusiasm.




Figure 7. Teacher AI implementation, feedback and confidence by self-reported policy familiarity

IV. What Is AI Doing to Learning, Teaching, and Assessment?

Do students learn better or worse with AI?

61% of students agree AI helps them learn effectively, the strongest positive consensus in the student data. At the same time, 23% agree that AI sometimes harms their understanding. **20% hold both positions at once: “AI helps my learning *and* it harms my understanding.”** We don’t think that this has to be a contradiction. The finding points at an ambivalence, as both positions can be true at the same time. AI is a tool and the outcome and its impacts depend on the way it is used. As mentioned earlier, students think about AI as a task-completer or as a thinking-partner, and we think that both can be true for the same individual. In interviews, one student articulates another explanation for the double position: “learn more, retain less.” AI gives access to more knowledge, but the consolidation that comes from effortful retrieval diminishes.

Beyond that, student concerns cluster into three distinct frameworks that can occur simultaneously in the same individual. The most common is **cognitive harm**: weakened critical thinking, reduced capacity to work independently, what many students call “brain rot.” Some individual students report discomfort, including guilt about over-reliance and self-recognition of habits they consider lazy. The second is **peer pressure to use AI**: AI raises the general standard of submitted work, so students now compete against AI-level output rather than peer-level output. Two students independently describe this mechanism. One says school has become more competitive because of AI and they do not like it. Both above and below average students use AI to boost results, which raises the bar for everyone and increases pressure to continue using AI. This happening more broadly risks creating a self-reinforcing cycle that is difficult for any individual school to break. The third framework is **social and existential**: job displacement, and long-term civilisational dependency on AI. This is present across schools but is rarely elaborated.

How does AI change the way teachers teach?

82% of teachers agree AI is a useful tool for student learning. 65% also agree AI hinders accurate assessment. Teachers are simultaneously enthusiastic about AI’s educational potential and alarmed about its institutional consequences, among others.

AI is a useful tool for student learning



AI hinders accurate student assessment



Strongly disagree Disagree Neutral Agree/Strongly agree

Figure 8. Teacher attitudes toward AI for learning and assessment

59% support AI integration in education; 18% oppose it. The benefits teachers identify most are: 1) differentiation of learning and class materials (especially in language teaching, where AI can adapt a single text across multiple proficiency levels), 2) AI as a private non-judgmental resource for students who lack confidence to ask questions publicly, 3) administrative tasks offloaded to AI allowing for more time devoted to “actual teaching.” When we follow up on the meaning of “actual teaching” a teacher states that they have more time to plan engaging classes and provide more detailed feedback.

In short, both groups perceive a tension in students’ use of AI in schools. For students the tension is internal: **AI helps me work but may be degrading my thinking.** For teachers it is institutional: **AI is valuable but is undermining the assessment system schools depend on.**

A subset of teachers has restructured practice in response: shifting assessment entirely to in-class or oral formats, capping or eliminating graded homework, prompting AI agents to provide directions rather than full answers, and redesigning tasks so they cannot be completed by AI alone. But these responses represent a minority. **The large neutral middle visible in the survey data has not yet shifted class design and assessment practice.** In some interviews we sense a hesitation because of deep uncertainty that a teacher sees coming from “90% danger and a promise of some future where it’s going to be beneficial that doesn’t seem to be here yet.”

What is happening to assessment?

Assessment practice is **shifting from evaluating outcomes toward evaluating processes:** oral and in-class formats, checkpoint models, and task designs that make AI use visible. One teacher described asking students to let a chatbot respond to the question and then requiring them to describe how they would have answered differently.

Respondents from all six schools independently converged on oral examination as their primary structural response to AI. School A reports near-zero false negatives across three years; however, teachers don’t primarily regard this as a success of AI use detection but rather as an alarming signal for the volume of AI work that remains unseen.

At IB-affiliated schools, out-of-class written work is no longer reliably attributable to individual students, and administrators we spoke with acknowledged this. Teachers sign academic honesty declarations certifying that student work is their own, even though they may not be able to fully verify it. One director of school demonstrated the problem by producing a Grade 7 IB internal assessment in approximately 20 minutes via five to seven prompts.

V. What Do Students and Teachers Actually Want?

Why are students using AI?

Teachers at four schools independently raised the same observation: **students who lean heavily on AI often do so not because they are lazy, but because they genuinely believe the tool produces better work than they can.** If these teachers are right, the real issue is a lack of confidence. Students are making a judgment about themselves, not just taking a shortcut.

Students themselves tell a somewhat different story. **The reasons they give are practical: AI is always there, it is fast, and it never makes them feel embarrassed.** A teacher requires effort to approach and will form an opinion of you. AI does not. Several students also mentioned that talking to a chatbot feels private in a way that talking to a person does not, though a smaller group said this privacy concern cuts the other way, making them reluctant to share anything personal with an AI at all.

For older students, especially those in the final years of the IB, the pressure of exams and a heavy workload adds another layer. These students describe AI as a quick fix when time runs short and stress runs high. The survey data reflect this: the class graduating in 2026, facing the most immediate exam pressure, reports the highest rate of heavy AI use (39%) while also expressing the greatest worry about what AI might be doing to their thinking (scoring 4.37 out of 5 on concern about cognitive harm, compared to 3.92 for other year groups).

What do students want from their schools?

When asked what they would change, students across all six schools gave a consistent answer: **teach us how to use this effectively and creatively.** But what they imagine that teaching would look like varies quite a bit. Some want practical instruction in how to write effective prompts. Others want to understand how AI actually works, what is happening inside the system. Some want frameworks for thinking about the risks, with one student comparing it to how schools approach drug prevention. One of the more technically fluent students in our dataset argued that AI should simply be woven into existing subjects rather than pulled out into its own course.

A dedicated AI course that came up in our interviews received mixed reviews. One student found it genuinely useful. Another said it covered things everyone already knew. Both assessments are probably accurate, just for different students.

What do teachers want from their schools?

Most of what teachers know about AI, they have learned from each other informally, in hallways, over lunch, in passing. **This peer exchange is valuable, but teachers at almost every school said it is not enough.** Only one school has a regular, structured session dedicated to sharing AI practice. What teachers describe wanting is not a one-day training or

a new software tool. It is an ongoing time to talk with colleagues, try things, and build confidence in a field that keeps changing beneath their feet.

At some schools, even this informal exchange is complicated by social dynamics. One teacher described feeling looked down on by colleagues for openly mentioning AI in their work. The resistance to AI among some staff is not just a matter of skepticism; it can carry a social edge. Teachers who have built their professional identity around deep subject knowledge can feel that AI undermines what they stand for, and that feeling can translate into real pressure on colleagues who engage with it openly.

Schools can make a clear case why students should work closely with their university advisor. A university advisor observed that more students are turning to AI tools to help write their application essays. It is obvious that chatting with AI is more convenient than working with advisors. However, AI lacks contextual understanding of the student, and can take away from their authentic expression. Establishing a long-term, collaborative relationship with advisors, sitting down to have deep conversations, is crucial for students to remain authentic and unique in their voice. School officials can explain the importance of working with university advisors early, and clearly demonstrate the downsides of using AI for application essays.

VI. Does School Policy Produce the Intended Outcomes?

What do students and teachers think about their AI policies?

Based on the interviews, almost all students agree with the AI rules across all six schools. However, some students were also frustrated that their school teaches only how *not* to use AI rather than how to use it for learning and problem solving. In the survey, 61% of students say they understand school AI rules, but only 48% say their teachers have explained the rules clearly. Self-perception bias may explain part of this 13% gap, though it could also be that students are learning about the rules outside of the classroom as well, through posters and gatherings.

Among teachers, 43% agree their school's AI policy is realistic to implement. But 46% selected “neither agree nor disagree,” the highest neutral rate on any teacher item. Combined with 36% neutral on “I implement policy consistently,” **a large teacher population is neither adopting nor resisting, likely due to uncertainty.** We think this silent middle is probably where the most change can happen.

Do students and teachers comply?

56% of students agree that they follow AI rules set by their teachers. **However, among these same students who said they follow the rules, 23% also say they have used AI in situations where an assignment explicitly discourages it.** One way to understand this is that students may still see themselves as rule-followers overall, even if they occasionally break the rules in practice. When asked in general terms, they tend to give a positive self-description. But when asked about specific situations, they are more likely to admit exceptions. Other explanations include a psychological effect because of the wording of the questions, as the first asked about a person (the teacher) and the second about a thing (the assignment). There could also be a survey bias due to participants agreeing by default.




Figure 9. Student self-reported rule compliance and rule-breaking by school

School B students report the highest rule compliance in the sample (mean 6.02 vs 5.23). School E students report the highest non-compliant use. While this may suggest a correlation between policy sophistication and compliance, our interview complicates the finding. One

student at School B stated that the traffic-light system “is not actually used” and that effective AI use is undetectable. Another described using AI constantly while not considering it rule-breaking because he does not submit AI output as a final product. It is possible that the most compliant-looking school in the survey data may have the widest gap between stated and actual behaviour.

Among teachers, 56% report implementing policy consistently. 62% say they give specific feedback when they suspect student AI use. In interviews, some teachers stated giving students a second chance and letting them resubmit the assignment. Others said that they ask students to explain the reasoning behind work, and would then regularly “get [their] suspicion proven to be true”.

Do teachers actually know how students are using AI?

75% of teachers believe students use AI instead of doing the work themselves, yet only 23% of students report themselves using AI when discouraged. The two groups are not answering the same question, so a direct comparison has limits, but the gap between these numbers suggests that teachers and students disagree sharply over how often students use AI when they are not supposed to. On one hand, teachers may be overly suspicious of students’ AI use; on the other hand, students may be underreporting their inappropriate AI use. Both could be true, though we are much more certain that students are underreporting their AI use.




Figure 10. AI visibility reported by teachers, % of agreeing or strongly agreeing

50% of teachers say students are not transparent about their AI use, and 48% believe that students use AI in ways impossible for teachers to detect. These attitudes show that **teachers mostly think AI use is hardly visible to them**. Nevertheless, teachers rarely rely on detection tools: 46% of teachers disagree with relying on them and only 24% agree. Multiple teachers voiced concerns about the effectiveness of AI detection tools because of the continuous improvement of AI and students getting more creative in using AI, like prompting AI to mimic their own writing level and proficiency. When teachers suspect students of using AI, 62% say they approach students directly for a conversation. The ways of approaching differ from letting a student resubmit, to giving the student a failing grade on the assignment.

In an attempt to make AI use more detectable, a common practice used by teachers is to require their students to complete their homework on documents where teachers see the history of the document. There, they can easily spot when a student takes AI generated output and pastes it into the document. **Students have found ways around this oversight: students at multiple schools described transcribing AI output by hand.** The trick describes an escalating cat and mouse problem. Some students frame undetectability as a skill and teachers increasingly regard it as an operational crisis requiring structural response.

On another note, **students and teachers independently bring up the risk of an AI feedback loop where AI-generated student work receives AI-generated teacher feedback**, with little human cognition engaged on either side. One teacher calls this dynamic “a vortex of pointlessness”. We do not know how widespread this is, but the fact that both sides of the classroom describe it independently suggests that it is not solely a hypothetical case.

VII. What Can Schools Control, and What Can't They?

What sits outside schools' direct control?

The International Baccalaureate program is the biggest structural influence on how schools think about and implement AI policies. Most administrators would like to see more engagement and thinking about curriculum adaptation from the IB. Currently, courseworks such as Extended Essays and Internal Assessments carry grades but cannot be easily verified, and many teachers are uncertain about how to navigate this situation.

Beyond the structural constraints of the IB, administrators navigate 3 external pressures that seem to outgrow institutional policy.

First is the speed at which AI is changing. Administrators at multiple schools described the “policy lag” where a policy can become obsolete before it is even fully implemented.

Second is the safeguarding risk regarding student mental health. Multiple teachers have voiced concerns about students giving too much credence to AI responses, because they treat AI as a human agent. In our study, we have found evidence of students substituting human contact with AI interactions for privacy and unlimited availability.

Third is a geopolitical infrastructure risk. Swiss international schools and their students remain almost entirely active on US-based platforms. As one teacher noted, the use of European-made models like Apertus is effectively zero. This creates a long-term dependency on foreign technology infrastructure that no individual school can mitigate.

What does the evidence suggest schools should prioritise?

1. Evaluating *process* instead of *outcome* seems to be a step towards academic integrity. Because teachers report an inability to detect AI use and most question the reliability of detection software, individual teachers are moving toward high-frequency, low-stakes verification. Teachers from all 6 schools have independently arrived at oral assessment and checkpoint models as the primary mechanism for establishing authorship. Completing this transition successfully and sustainably, however, requires institutional action rather than individual effort.

2. Teacher practice has a *higher* impact than policy documents. The data shows that a total absence of a policy framework is associated with worse outcomes. However, solely focusing on making the policy more comprehensive does not seem to increase student compliance. Making sure that teachers are confident about implementing the current policy is much more impactful than making the policy document more sophisticated. We see the policy document as the guiding lighthouse and teachers as the ferryperson that accompanies and leads the student along their journey.




Figure 11. Student rule compliance and teacher policy implementation consistency by school [self-reported]

3. Addressing the decrease in self-confidence. As noted before, teachers at four schools independently pointed to low self-confidence as the main reason students over-rely on AI. Students themselves more often cite convenience and accessibility. We cannot say which explanation is more accurate, and both may be true for different students or in the same individual concurrently. But if the “lack of confidence” explanation is right, it would suggest that the most effective response is not stricter discouragement of AI use, but **finding ways to rebuild students’ belief in their own intellectual capacity.**

4. What students and teachers want. Both students and teachers consistently emphasized guiding students on how to use AI, rather than trying to ban it. Students want formal AI education even though administrators are universally sceptical of AI-specific skills as durable additions to the curriculum. We agree that models and cutting-edge applications often transform too quickly to create meaningful and sustainable coursework that focuses on teaching platform-specific skills. Therefore, we believe that teaching students the working principles and discussing the multifaceted impacts of the technology may be most helpful for students.

Teachers expressed a clear demand for structured dialogue to replace the current reliance on informal peer exchange. Right now, most learning happens through ad hoc teacher conversations, which are uneven, opaque, and hard to scale. Our data suggest that the most effective professional development would formalise these exchanges into regular, structured spaces where teachers can openly share practices, concerns, and failures.

5. AI may be clarifying the role of educators and the purpose of education. In interviews, teachers describe a reassessment of what their role is for. They ask themselves: what does

school provide when AI can deliver knowledge on demand? The answers we hear, articulated in different ways at different schools, point in a similar direction: accompaniment, the discipline of thinking, and the social dimensions of learning. As one teacher puts it: “The school doesn’t have the copyright on knowledge anymore. So if pedagogy [from greek “paidagōgos”] means walking next to someone, students still need people to walk with.”

Three metaphors for the teacher’s role emerged independently across the schools we visited. Some frame the teacher as a **foundation builder**, someone who ensures students have the literacies and judgment needed to make AI use meaningful. A second metaphor is the **mountain guide**: the teacher does not simply point out the destination, but leads the route, manages group dynamics, and helps students move safely through difficult terrain. The third is the **values anchor**, where the teacher is not only a source of support but also the steady human presence that holds the school’s values in place through daily interactions and relationships.

A school that sees teachers as foundation builders may invest in curriculum and assessment redesign. A school that sees teachers as mountain guides may invest in classroom structure and AI integration tools. A school that sees teachers as values anchors may invest in community and school culture. Taken together, these framings point to a shared conclusion: AI may be changing how education is delivered, but it has not changed what education is for. What AI is asking schools to do is rethink the educator’s role, and in doing so, clarify which parts of education are technical, which are relational, and which are grounded in unchanging values.

Appendix A: Interview Questions

While we used the following questions to guide the interviews, the interviews were not rigidly structured through them. Different interviews focused on different aspects of the questions because of factors including but not limited to interviewee comfort and knowledge, interview length, and interview setting.

Guiding Questions For Administrator Interviews

1. Meaning of AI Use
 - a. When people at this school talk about “using AI,” what kinds of activities do you think they usually mean?
 - b. Are there uses that feel clearly acceptable versus clearly problematic?
 - c. Are there gray areas that come up repeatedly?
2. Policy Framing
 - a. Does the school currently have a formal policy or guidelines regarding AI use?
 - b. What were the main considerations behind how that policy was designed?
 - c. Was it driven more by pedagogy, assessment concerns, risk management, or external expectations?
 - d. Was there internal debate or disagreement during its development?
 - e. Do you see it as a stable policy, or something still evolving?
3. In-the-Moment Decision-Making
 - a. When you think about teachers or students deciding whether to use AI in a specific situation, what factors do you think weigh most heavily for them?
 - b. Time pressure? Difficulty? Stakes of the task?
 - c. Do you think these decisions are more rule-based or judgment-based in practice?
4. Emotional Orientation & Perceived Pressure
 - a. How would you describe the general emotional climate around AI at the school?
 - b. Do people seem more curious, anxious, resistant, or resigned?
 - c. Do you sense pressure — for students to keep up, or for teachers to adapt?
 - d. Does AI use feel optional, expected, or unavoidable?
5. Perceived Norms vs Official Rules
 - a. How closely do you think everyday AI practices align with the school’s formal expectations?
 - b. Where do you see the biggest gaps between policy and practice?
 - c. Are there areas where ambiguity is intentional or unavoidable?
 - d. How transparent do you think people feel they can be about AI use?
6. Trust, Authority, and Substitution
 - a. In what ways do you think AI is changing how authority or trust operates in the school — between students, teachers, and the institution?
 - b. Are there situations where AI is treated as more reliable than human guidance?
 - c. Do you think AI complicates how learning or understanding is recognized?

- d. Have you noticed any shifts in trust — epistemic or relational — as AI use increases?
- 7. Learning, Authorship, and Ownership
 - a. How does the school currently think about learning and authorship in an AI-supported environment?
 - b. Is struggle still seen as essential to learning?
 - c. How do you distinguish assistance from substitution?
 - d. Are these distinctions clear to students and teachers?
- 8. Moral Comfort, Risk, and Justification
 - a. Where do you think people feel most morally uncertain or uncomfortable about AI use?
 - b. Are there practices people quietly tolerate but wouldn't openly endorse?
 - c. What risks concern you most — academic, relational, reputational?
- 9. Desired Futures and Direction
 - a. Looking ahead, what do you think the school should be prioritizing in how it approaches AI?
 - b. What new literacies do students need to avoid becoming overly dependent or passive?
 - c. If you could change one thing about how AI is currently handled, what would it be?
 - d. What do you hope students take away from learning with AI?

Guiding Questions For Teacher Interviews

- 1. Conceptualizing Use
 - a. When you think about student AI use, what behaviors/tools come to mind?
 - b. What uses feel supportive? What feels like substitution?
 - c. What are the hardest gray areas of student use?
- 2. Policy Framing
 - a. What are the guidelines or the policy around AI that has been communicated to you?
 - b. Is it clear and usable in daily teaching?
 - c. Does the policy reflect classroom reality? Are you forced to adapt or interpret it?
- 3. Decision-Making
 - a. What do you think drives student AI decisions?
 - b. Do you think students usually think in terms of “rules,” or in terms of what feels reasonable?
 - c. How do you decide when to allow, discourage, or ignore AI use?
- 4. Emotional Orientation
 - a. Do you feel pressure to adapt or integrate AI?
 - b. Do you sense competition or comparison among students because of AI?
 - c. Does AI feel like an opportunity, a burden, or something in between?

5. Trust & Authority
 - a. Has widespread AI use changed how you assess student understanding or authentic voice?
 - b. Can you usually distinguish between a student using AI to understand a concept versus using it to bypass the work?
 - c. Does knowing students may use AI affect how much you trust what you see in their work?
 - d. Has it changed how you interact with students or give feedback?
6. Learning & Authorship
 - a. How do you currently think about learning and authorship when AI is involved?
 - b. Are these expectations clear to students?
7. Moral Comfort
 - a. Are there aspects of AI use in teaching or assessment that make you feel uncertain or uneasy?
 - b. What risks of AI concern you the most?
8. Desired Futures
 - a. Looking ahead, what do you think would help teachers and students use AI in more constructive ways?
 - b. What skills or literacies should be emphasized?
 - c. What kind of guidance or support would be most useful?

Guiding Questions For Student Interviews

1. Use & Learning
 - a. Walk me through the last time you used AI?
 - b. How did you learn to use it?
 - c. What were you working on? Why did you decide to use AI? Did you hesitate?
2. Learning Style with AI
 - a. What do you do when you don't understand an AI answer?
 - b. How often do you fact-check AI output?
 - c. When you don't understand an answer by AI, what do you do? Ask to explain? Check other sources? Use as-is?
3. Policy
 - a. What rules and guidelines around AI in school were given to you and by whom?
 - b. Were these consistent across teachers?
 - c. Do they make sense to you?
4. Psychological Safety & Norms
 - a. When you have a question or need help on school work, are you more comfortable asking AI, a teacher, a friend or someone else? Why?
 - b. Talk openly about AI with friends/teachers?

- c. Does AI change how you learn or ask questions?
 - d. Does AI make classes/certain subjects more/less valuable?
- 5. Education Values
 - a. What are your feelings about how you use AI?
 - b. Do you think AI is beneficial to actual learning?
- 6. Future Directions
 - a. How do you feel about the technology? Do you see any particular risks?
 - b. If there was one thing that you could change about how your school approaches AI, what would it be?
 - c. Would you want to be taught about AI? How and by whom?
 - d. What would happen if AI stopped existing right now?

Appendix B: Notes

Full Survey Data. The researchers have visualized the full survey data into an interactive dashboard in order to ensure anonymity of the respondents and to facilitate understanding of the dataset.

AI Use for This Report. The researchers used AI, specifically Claude, to format the report and to create the chart visualizations.